

# Homework3

Yuning Li

2024-02-21

```
library('splines')      ## for 'bs'
library('dplyr')        ## for 'select', 'filter', and others

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
library('magrittr')     ## for '%>%' operator
library('glmnet')       ## for 'glmnet'

## Loading required package: Matrix
## Loaded glmnet 4.1-8
### Linear regression examples ###

## load prostate data
prostate <-
  read.table(url(
    'https://web.stanford.edu/~hastie/ElemStatLearn/datasets/prostate.data'))

## split prostate into testing and training subsets
prostate_train <- prostate %>%
  filter(train == TRUE) %>%
  select(-train)

prostate_test <- prostate %>%
  filter(train == FALSE) %>%
  select(-train)

## predict lpsa consider all other predictors
fit <- lm(lpsa ~ ., data=prostate_train)

## functions to compute testing/training error
L2_loss <- function(y, yhat)
  (y-yhat)^2
```

```

error <- function(dat, fit, loss=L2_loss)
  mean(loss(dat$lcavol, predict(fit, newdata=dat)))

## train_error
error(prostate_train, fit)

## [1] 1.647994

## testing error
error(prostate_test, fit)

## [1] 1.474148

## use glmnet to fit ridge regression
form <- lpsa ~ lweight + age + lbph + lcp + pgg45 + lcavol + svi + gleason
x_inp <- model.matrix(form, data=prostate_train)
y_out <- prostate_train$lpsa
# Use glmnet to fit ridge regression
fit_ridge <- glmnet(x = x_inp, y = y_out, alpha = 0, lambda = seq(0.5, 0, -0.05))

# Display the coefficients for different lambda values
print(fit_ridge$beta)

## 9 x 11 sparse Matrix of class "dgCMatrix"
##   [[ suppressing 11 column names 's0', 's1', 's2' ... ]]
##
## (Intercept)      .      .      .      .      .
## lweight      0.515697286  0.525400575  0.5354170019  0.545775814  0.556311619
## age          -0.004889276 -0.005637483 -0.0064660393 -0.007390172 -0.008436051
## lbph          0.111264821  0.113946024  0.1167762533  0.119749176  0.122883708
## lcp           0.015168605  0.007943168 -0.0003765233 -0.010208916 -0.022217776
## pgg45         0.004390331  0.004515174  0.0046637167  0.004840594  0.005062742
## lcavol        0.329947512  0.341656357  0.3545916725  0.369018247  0.385454679
## svi           0.541710372  0.553480497  0.5659053246  0.579302886  0.594067995
## gleason       0.063402833  0.061373728  0.0588793776  0.055915056  0.052224299
##
## (Intercept)      .      .      .      .      .
## lweight      0.567083034  0.577904675  0.588577079  0.598728953  0.607640794
## age          -0.009617693 -0.010968049 -0.012524059 -0.014335139 -0.016468427
## lbph          0.126178477  0.129642591  0.133272193  0.137050755  0.140935937
## lcp          -0.036796539 -0.054965954 -0.078076538 -0.108262379 -0.149116876
## pgg45         0.005340520  0.005702607  0.006189804  0.006872589  0.007882003
## lcavol        0.404236189  0.426145259  0.452197566  0.483962606  0.524027412
## svi           0.610303587  0.628512630  0.649279407  0.673464940  0.702375090
## gleason       0.047591554  0.041501934  0.033149650  0.021047902  0.002236316
##
## (Intercept)      .
## lweight      0.613980850
## age          -0.019008894
## lbph          0.144825646
## lcp          -0.206830251
## pgg45         0.009471377
## lcavol        0.576765946
## svi           0.737805454
## gleason      -0.029463345

```

```

# Tune the value of lambda based on cross-validation
cv_fit_ridge <- cv.glmnet(x = x_inp, y = y_out, alpha = 0)

# Display the optimal lambda
best_lambda <- cv_fit_ridge$lambda.min
print(best_lambda)

## [1] 0.08788804

## functions to compute testing/training error with glmnet
error <- function(dat, fit_ridge, lam, form, loss=L2_loss) {
  x_inp <- model.matrix(form, data=dat)
  y_out <- dat$lpsa
  y_hat <- predict(fit_ridge, newx=x_inp, s=lam) ## see predict.elnet
  mean(loss(y_out, y_hat))
}

## train_error at lambda=0
error(prostate_train, fit_ridge, lam=0, form=form)

## [1] 0.4391999

## testing error at lambda=0
error(prostate_test, fit_ridge, lam=0, form=form)

## [1] 0.5214441

## train_error at lambda=0.03
error(prostate_train, fit_ridge, lam=0.05, form=form)

## [1] 0.4424079


## testing error at lambda=0.03
error(prostate_test, fit_ridge, lam=0.05, form=form)

## [1] 0.5019552

## compute training and testing errors as function of lambda
err_train_1 <- sapply(fit_ridge$lambda, function(lam)
  error(prostate_train, fit_ridge, lam, form))
err_test_1 <- sapply(fit_ridge$lambda, function(lam)
  error(prostate_test, fit_ridge, lam, form))

## plot test/train error
plot(x=range(fit_ridge$lambda),
     y=range(c(err_train_1, err_test_1)),
     xlim=rev(range(fit_ridge$lambda)),
     type='n',
     xlab=expression(lambda),
     ylab='train/test error')
points(fit_ridge$lambda, err_train_1, pch=19, type='b', col='darkblue')
points(fit_ridge$lambda, err_test_1, pch=19, type='b', col='darkred')
legend('topright', c('train', 'test'), lty=1, pch=19,
      col=c('darkblue', 'darkred'), bty='n')

```



```
colnames(fit_ridge$beta) <- paste('lam =', fit_ridge$lambda)
print(fit_ridge$beta %>% as.matrix)
```

```
##          lam = 0.5   lam = 0.45   lam = 0.4   lam = 0.35   lam = 0.3
## (Intercept) 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
## lweight    0.515697286 0.525400575 0.5354170019 0.545775814 0.556311619
## age        -0.004889276 -0.005637483 -0.0064660393 -0.007390172 -0.008436051
## lbph        0.111264821 0.113946024 0.1167762533 0.119749176 0.122883708
## lcp         0.015168605 0.007943168 -0.0003765233 -0.010208916 -0.022217776
## pgg45       0.004390331 0.004515174 0.0046637167 0.004840594 0.005062742
## lcavol      0.329947512 0.341656357 0.3545916725 0.369018247 0.385454679
## svi         0.541710372 0.553480497 0.5659053246 0.579302886 0.594067995
## gleason     0.063402833 0.061373728 0.0588793776 0.055915056 0.052224299
##          lam = 0.25   lam = 0.2   lam = 0.15   lam = 0.1   lam = 0.05
## (Intercept) 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
## lweight    0.567083034 0.577904675 0.588577079 0.598728953 0.607640794
## age        -0.009617693 -0.010968049 -0.012524059 -0.014335139 -0.016468427
## lbph        0.126178477 0.129642591 0.133272193 0.137050755 0.140935937
## lcp        -0.036796539 -0.054965954 -0.078076538 -0.108262379 -0.149116876
## pgg45       0.005340520 0.005702607 0.006189804 0.006872589 0.007882003
## lcavol      0.404236189 0.426145259 0.452197566 0.483962606 0.524027412
## svi         0.610303587 0.628512630 0.649279407 0.673464940 0.702375090
## gleason     0.047591554 0.041501934 0.033149650 0.021047902 0.002236316
##          lam = 0
## (Intercept) 0.000000000
## lweight    0.613980850
## age        -0.019008894
## lbph        0.144825646
## lcp        -0.206830251
## pgg45       0.009471377
## lcavol      0.576765946
## svi         0.737805454
## gleason    -0.029463345
```

